

CAVA v1.2.3

documentation

CONTENTS

1	INTRODUCTION	-----	2
2	INSTALLATION	-----	2
3	RUNNING CAVA	-----	3
4	CONFIGURATION FILE	-----	4
5	INPUT FILE	-----	6
6	REFERENCE GENOME	-----	6
7	TRANSCRIPT DATABASE	-----	6
8	DBSNP DATABASE	-----	7
9	VARIANT ANNOTATION	-----	7
10	THE CLASS ONTOLOGY	-----	11
11	OUTPUT	-----	14
12	VARIANT CALL FILTERING	-----	17
13	THE ENSEMBL_DB TOOL	-----	18
14	DEFAULT TRANSCRIPT DATABASE	---	20
15	THE DBSNP_DB SCRIPT	-----	20
16	RUN TIMES	-----	22
17	LOG FILE	-----	22
18	EXAMPLES	-----	22
19	CONTACT	-----	29
APPENDIX 1			-----30
APPENDIX 2			-----38

1 INTRODUCTION

CAVA (Clinical Annotation of Variants) is a lightweight, flexible, fast and easy-to-use Next-Generation Sequencing (NGS) variant annotation tool specializing in transcript-level annotation. This detailed documentation describes all features of CAVA and its accompanying transcript database preparation tool, `ensembl_db`.

After first introducing the CAVA user interface and discussing its functionalities, some usage examples are shown in the second part of this documentation.

2 INSTALLATION

Dependencies:

To install and run CAVA v1.2.3 you will need the following dependencies installed:

- Python 2.7.9 or later (Python2 series)
- GCC and GNU make
- virtualenv

If your system is missing GCC and GNU make, these can be installed as follows:

- On a Mac, the easiest way to set them up is to install Xcode Command Line Tools.
- On a Debian or Ubuntu Linux, they can be set up by installing the build-essential package:

`sudo apt-get install build-essential`

If not already installed on your system, virtualenv can be set up by:

`pip install virtualenv`

Installation on Linux or Mac:

CAVA v1.2.3 can be downloaded from

<https://github.com/RahmanTeam/CAVA/releases/tag/v1.2.3>

in either .zip or .tar.gz format. To unpack these run one of the following commands:

`unzip CAVA-1.2.3.zip`

or:

`tar -xvzf CAVA-1.2.3.tar.gz`

and then you can install CAVA with the following command from the CAVA-1.2.3 directory:

`./install.sh`

CAVA uses virtualenv and pip to manage all its extra dependencies, which means that it will not clutter up your system by installing things globally. Everything it installs will go into a sub-directory in the CAVA-1.2.3 directory. If you delete CAVA then everything it has installed will also be deleted. Once the installation script has finished successfully, CAVA is ready for use.

The .zip or .tar.gz file also includes the transcript database preparation tool, `ensembl_db` (see Section 13), a template configuration file, and a default whole exome transcript database (Ensembl release 75).

3 RUNNING CAVA

Once downloaded and correctly installed, CAVA can be run with the following simple command:

`CAVA-1.2.3/cava -c config.txt -i input.vcf -o output`

CAVA requires three command line arguments: the name of the configuration file (-c), the name of the input file (-i) and the prefix of the output file name (-o).

The configuration file should contain the user-specified settings (see Section 4). The input file should provide the set of variant calls to be annotated. This can be in standard VCF format or a TXT file (see below for possible input formats). Finally, the output prefix is used for the name of the output file and the log file.

Optionally, by using option -s, the annotated variant set will be written to the standard output (stdout) rather than to an output file. A log file will still be created, if required.

Optionally, CAVA can be run with multithreading using option -t (see Section 3.1).

Help information for the command line arguments can be requested using the -h argument:

`CAVA-1.2.3/cava -h`

3.1 Multithreading

CAVA has a built-in multithreading feature which allows variant annotation to be performed in multiple parallel processes, resulting in significant speedup. Each process generates a temporary output file, which is automatically merged at the end to create the final output. To

use multithreading, one only needs to apply command line option flag `-t` and specify the number of processes.

For example, using 4 processes:

CAVA-1.2.3/cava -c config.txt -i input.vcf -o output -t 4

(See Section 16 for an example of how multithreading speeds up CAVA.)

3.2 Default configuration file path

If one is repeatedly using the same configuration file, it is possible to specify it as a ‘default’ configuration file by adding the path pointing to the configuration file into the `default_config_path` file in the CAVA directory. The `default_config_path` file should contain a single line, e.g. `/path/to/config.txt`.

This will simplify the command for running CAVA:

CAVA-1.2.3/cava -i input.vcf -o output

Running the above command with no `-c` option, CAVA will use the configuration file defined in the `default_config_path` file.

4 CONFIGURATION FILE

As mentioned above, the configuration file should contain all user-specified options. Here we describe the list of possible option flags (see more detailed discussions in next sections). Example configuration files are also given in Section 18 and a template is included in the `.zip` or `.tar.gz` package.

Each option set in the configuration file should be given in the format of

`@flagname = value`

CAVA understands the following option flags (some are mandatory and some are optional with the default values shown below):

@inputformat : Input file format. (Optional. Possible values: VCF or TXT. Default value: VCF)

@outputformat : Output file format. (Optional. Possible values: VCF or TSV. Default value: VCF)

@reference : Absolute path to reference genome file. (Mandatory)

@ensembl : Absolute path to Ensembl transcript database file. (Optional. If not given, the default transcript database will be used – see Section 14)

@dbSNP : Absolute path to dbSNP database file. (Optional. If not given, no SNP-based annotation will be performed)

@nonannot : Boolean flag specifying if variants which received neither transcript nor dbSNP annotations are to be included in the output. (Optional. Possible values: TRUE or FALSE. Default value: TRUE)

@filter : Boolean flag specifying if only records with PASS filter value are included in the output. (Optional. Possible values: TRUE or FALSE. Default value: FALSE)

@type : Types of variants to be annotated. (Optional. Possible values: ALL, SUBSTITUTION, INDEL, INSERTION, DELETION or COMPLEX. Default value: ALL)

@target : Name of compressed BED file specifying genomic regions variant annotation is restricted to. (Optional)

@genelist : Name of file providing a list of the gene identifiers variant annotation is restricted to. Gene identifiers are to be given on separate lines in the file. (Optional)

@transcriptlist : Name of file providing a list of the transcript identifiers variant annotation is restricted to. Transcript identifiers are to be given on separate lines in the file. (Optional)

@snplist : Name of file providing a list of the dbSNP identifiers variant annotation is restricted to. dbSNP identifiers are to be given on separate lines in the file. (Optional)

@logfile : Boolean flag specifying if log file is written. (Optional. Possible values: TRUE or FALSE. Default value: FALSE)

@ontology : Which ontology is used for reporting functional class assignment. (Optional. Possible values: CLASS, SO or BOTH. Default value: BOTH)

@impactdef : Definition of variant impact levels (reported by the IMPACT annotation flag)

@givealt : Boolean flag specifying if alternative most 5' and CLASS/SO annotations are to be outputted. (Optional. Possible values: TRUE or FALSE. Default value: TRUE)

@ssrange : Number of bases into the intron defined as the splice site region (as used in the CLASS flag). (Optional. Possible values: integer ≥ 6 . Default value: 8)

@prefix : Boolean flag specifying if annotation flag names start with the prefix 'CAVA_' in VCF output (Optional. Possible values: TRUE or FALSE. Default value: FALSE)

5 INPUT FILE

The input file (defined by command line argument -i) contains all variant calls to be annotated. It may follow two formats, VCF or TXT (specified in the configuration file by option @inputformat).

5.1 Input in VCF format

If VCF format is used, the input file should follow the standard [VCF 4.1 specification](#). A single VCF record can describe a multiallelic variant which will be considered by CAVA as multiple different variant calls. VCF records may also include genotype call information.

5.2 Input in TXT format

If TXT format is used, the input file must be written in the following tab-delimited 5-column format describing the variant ID, chromosome name, genomic position, reference allele and alternative allele. If a variant is multiallelic, the alternative alleles must be comma-separated:

#ID	CHROM	POS	REF	ALT
1	13	1324552	G	A,TC
...				

Note that this format does not contain information about the filter value, therefore filter=PASS will be set for each variant.

6 REFERENCE GENOME

In order to annotate variants, CAVA requires a reference genome sequence. The reference genome (specified in the configuration file by option @reference) must be a FASTA file indexed with *samtools faidx*. Both the FASTA file and the index file (e.g. human_g1k_v37.fasta and human_g1k_v37.fasta.fai) should be available in the same directory. Chromosome names in the FASTA file can either be of the format '14' or 'chr14'.

For example, the GRCh37 reference genome sequence (together with the corresponding .fai file) can be downloaded from the 1000G website:

- ftp://ftp.1000genomes.ebi.ac.uk/vol1/ftp/technical/reference/human_g1k_v37.fasta.gz
- ftp://ftp.1000genomes.ebi.ac.uk/vol1/ftp/technical/reference/human_g1k_v37.fasta.fai

7 TRANSCRIPT DATABASE

To perform transcript-based annotation, CAVA requires a local transcript database that contains all relevant information about each transcript (strandedness, genomic positions of

exons, etc.). The transcript database file needs to have a special format, compressed with bgzip and indexed with Tabix. A simple tool (`ensembl_db`) is provided to generate the correct, compressed and Tabix-indexed transcript database file based on a user-specified Ensembl release (see Section 13 for a detailed description of the `ensembl_db` tool).

CAVA v1.2.3 also contains a default transcript database (see Section 14), based on Ensembl release 75, that is ready for use after installation.

When using a non-default transcript database, the name of the transcript database file must be specified in the configuration file by the option `@ensembl`.

8 DBSNP DATABASE

In order to perform SNP annotation, CAVA requires a local dbSNP database that contains information about each SNP (i.e. its identifier, genomic position and alt allele). The dbSNP database file needs to have a special format compressed with bgzip and indexed with Tabix. As with the transcript database discussed above, a simple tool, `dbsnp_db`, is included that can be used to generate the compressed and Tabix-indexed dbSNP database file (see Section 15 for a detailed description of the `dbsnp_db` tool).

The name of the dbSNP database file is specified in the configuration file by option `@dbsnp`.

9 VARIANT ANNOTATION

CAVA outputs information about variant calls in various annotation flags. Two types of variant annotation are supported: transcript-based and SNP-based. In transcript-based annotation, variant calls are annotated based on the information stored in the local Ensembl transcript database file. In SNP-based annotation, variants are annotated based on the local dbSNP database file.

This section discusses all components of variant annotation. Section 11 describes the exact way these components are written into different output formats.

One annotation flag which is always outputted (regardless of performing transcript-based or SNP-based annotation) is 'TYPE' which refers to the type of variant call (e.g. TYPE=SUBSTITUTION). This annotation flag has four possible values: SUBSTITUTION, INSERTION, DELETION and COMPLEX, referring to base substitutions, insertions, deletions and complex indels, respectively.

9.1 Transcript-based annotation

If option flag `@ensembl` is set in the configuration file, CAVA will search the transcript database for each variant call to find Ensembl transcripts that overlap with the variant. Note

that CAVA allows overlapping transcripts and a variant may overlap with multiple Ensembl transcripts (see Section 11.1.2).

9.1.1 The TRANSCRIPT, GENE, GENEID and TRINFO annotation flags

Variant calls overlapping with at least one transcript are annotated with the transcript identifiers in the annotation flag 'TRANSCRIPT' (e.g. TRANSCRIPT=ENST00000380152).

Transcript-overlapping variants are also annotated with the corresponding gene names (i.e. HGNC symbols) in the annotation flag 'GENE' (e.g. GENE=BRCA2) and the gene identifier in the annotation flag 'GENEID' (e.g. GENEID=ENSG00000139618). Basic information about the transcripts (strandedness, length of transcript, number of exons, length of coding DNA + UTR and length of protein product, respectively) are added in the 'TRINFO' annotation flag (e.g. TRINFO= +/84.8kb/27/12.0kb/3418). For example, a TRINFO annotation flag value of '+/84.8kb/27/12.0kb/3418' means that the transcript is forward stranded (+), with a total sequence length of 84.8 kb, including 27 exons which together cover 12.0 kb of coding region + UTR and correspond to a 3418 aa protein sequence (excluding the termination codon).

9.1.2 The LOC annotation flag

The location of the variant within the transcript is added in the annotation flag 'LOC' (e.g. LOC=Ex8). As exemplified below, the LOC flag may refer to an exon, an intron, a UTR3 or a UTR5 region.

Alternatively, if a variant overlaps the boundary between two different regions (e.g. exon and intron), this is referred to as 'LOC=L1-L2', where L1 and L2 are the locations of the two endpoints of the variant.

Examples:

Variant within exon 12: LOC=Ex12

Variant within the intron separating exons 5 and 6: LOC=In5/6

Variant within 5' untranslated region: LOC=UTR5

Variant within 3' untranslated region: LOC=UTR3

Variant overlapping the 3' boundary of exon 7: LOC=Ex7-In7/8

Variant overlapping exon 28 and the 3' untranslated region: LOC=Ex28-UTR3

Note that some Ensembl transcripts contain artificial introns of very short length (1,2,4 or 5 bp) called 'frameshift introns', added by the Ensembl genebuild. Since these may indicate errors in either the reference genome assembly or in the aligned cDNA, variants overlapping frameshift introns may actually be exonic or not real. CAVA thus flags variants affecting these introns, reporting this information in the LOC annotation flag, e.g. LOC=fsIn8/9 is a reference to the frameshift intron separating exons 8 and 9.

9.1.3 The CSN, CLASS, SO, IMPACT and ALTFLAG annotation flags

The CSN flag:

The Clinical Sequencing Notation (CSN) v1.0 is used for clinical variant annotation. It is based on the Human Genome Variation Society (HGVS) with minor amendments to allow high-volume automated outputs from NGS pipelines. The CSN is described in detail in Appendix 1.

The CSN annotation of a variant affecting a transcript is outputted in the annotation flag 'CSN' (e.g. CSN=c.421A>G_p.Thr141Ala). For indels, CSN is created based on the correctly aligned (strand-aware) representation of the variant. Note that the CSN annotation is given both on DNA (c.) and protein (p.) level if appropriate (otherwise only c. is given).

The CLASS and SO flags:

Variants are also classified according to a simple ontology and the class of variant is outputted in the annotation flag 'CLASS' (e.g. CLASS=SY). See Section 10 for detailed description of the CLASS ontology and the definition of different classes. Alternatively, one can output the Sequence Ontology (SO) annotation instead of the CLASS ontology or both ontologies can be reported. One can use the @ontology flag in the configuration file to control which ontologies are outputted. The @ontology flag has three possible values: CLASS, SO or BOTH, specifying that the CLASS ontology, the SO ontology or both are reported, respectively. The default value is BOTH, thus by default both CLASS and SO are outputted. For the description of the SO classification, please refer to the Sequence Ontology website: <http://www.sequenceontology.org/>. For a comparison of the CLASS and SO ontologies, see the table in Section 10.

The IMPACT flag:

CAVA can also stratify variants into groups of likely similar impact based on their CLASS annotations. This information is reported in the IMPACT annotation flag. The value of the IMPACT flag is an integer indicating decreasing levels of impacts, with the following default corresponding CLASS values:

IMPACT	CLASS
1	ESS, FS, SG
2	NSY, SS5, IF, IM, SL, EE
3	SY, SS, INT, 5PU, 3PU

Furthermore, one can define a custom mapping from CLASS terms to IMPACT values by using the @impactdef option flag in the configuration file. Different impact levels are separated by | and a comma-separated list of CLASS terms must be given for each level. For example, the following line defines the default mapping described above:

@impactdef=ESS,FS,SG | NSY, SS5,IF,IM,SL,EE | SY,SS,3PU,5PU,INT

Note that the IMPACT annotation flag is always outputted unless “@impactdef=.” is set in the configuration file. The IMPACT flag is reported even if the CLASS flag is not outputted (i.e. @ontology=SO).

The ALTFLAG flag:

Some indels have alternative representations which may even change their CLASS or SO annotations. CAVA recognizes variants with alternative annotations and outputs this information in the annotation flag 'ALTFLAG' (e.g. ALTFLAG=None). If, for example, only the CLASS ontology is reported, the ALTFLAG flag has three possible values: 'None', 'AnnNotClass' and 'AnnAndClass'. If the variant has the same annotation regardless of its left or right alignment, ALTFLAG=None is given. If an indel has an alternative annotation but the same CLASS annotation for both representations, ALTFLAG=AnnNotClass is given. Finally, if the indel has different CLASS annotations depending on its alignment, ALTFLAG=AnnAndClass is outputted, referring to the fact that the different representations may be interpreted as having different functional consequences.

Similarly, when only the SO ontology is reported, the ALTFLAG annotation flag has the following three possible values: 'None', 'AnnNotSO' and 'AnnAndSO'.

If both CLASS and SO ontologies are reported, the possible values of ALTFLAG are: 'None', 'AnnNotClassNotSO', 'AnnAndClassNotSO', 'AnnAndSONotClass', 'AnnAndClassAndSO', indicating whether the different indel representations would result in the same annotation and same CLASS and/or SO annotations.

9.1.4 The PROTPOS, PROTREF and PROTALT annotation flags

In addition to the CSN annotation, CAVA outputs the affected protein position, reference protein amino acid(s) and alternate amino acid(s) of each protein-altering variant. The PROTPOS annotation flag reports the coordinates of amino acids changed by the variant (e.g. PROTPOS=312 or PROTPOS=45-48). The annotation flag PROTREF provides the reference amino acid sequence (e.g. PROTREF=GCHAR), while the annotation flag PROTALT gives the alternate amino acid sequence (e.g. PROTALT=RTY). The PROTREF and PROTALT flags report the amino acid sequence according to the single letter amino acid code.

For frameshifting indels, the PROTPOS and PROTREF flags report the position of the first affected amino acid and the first affected amino acid, respectively, while PROTALT is set to empty (PROTALT = .). For variants not affecting the protein sequence (e.g. intronic or UTR variants), PROTPOS, PROTREF and PROTALT are all set to empty (e.g. PROTPOS = .).

9.1.5 Outputting alternative annotations: ALTANN and ALTCLASS/ALTSO

If the @givealt option flag is set to TRUE in the configuration file, CAVA will also output the most 5' alternative sequence and CLASS and/or SO annotations instead of just indicating that alternative annotations exist in the ALTFLAG flag. If the @givealt option is switched on, additional annotation flags (ALTANN and ALTCLASS and/or ALTSO) are given for each

variant, providing the alternative sequence and CLASS and/or SO annotations (if any), respectively. While the correct CSN annotation is reported by the CSN annotation flag, if relevant, the most 5' effect on the protein is outputted in the ALTANN annotation. Note that if @givealt is used, the ALTFLAG annotation flag described above is not reported.

9.1.6 Variants located outside transcripts

The transcript-based annotations described above are only given if both endpoints of the variant are located within the transcript. If there is only partial overlap between the variant and the transcript (e.g. the starting position of a deletion lies within the transcript but the ending position lies outside), the partial overlap is reported in the following way:

The TRANSCRIPT, GENE, GENEID and TRINFO flags describing the transcript are given as above, however LOC=OUT is reported and empty CSN, PROTPOS, PROTREF, PROTALT, CLASS, SO, ALTFLAG, ALTANN, ALTCLASS and ALTISO flag values are outputted (e.g. CSN=.).

If no transcripts are found in the transcript database with which the variant call overlaps, all transcript-specific annotation flags (TRANSCRIPT, GENE, GENEID, TRINFO, LOC, CSN, PROTPOS, PROTREF, PROTALT, CLASS, SO, ALTFLAG, ALTANN, ALTCLASS and ALTISO) will have empty value (e.g. TRANSCRIPT=.).

9.2 SNP annotation

If option flag @dbSNP is set in the configuration file, for each base substitution CAVA will search the dbSNP database file to find registered SNPs. If a variant is identified as a known SNP, the annotation flag 'DBSNP' will output its dbSNP identifier (e.g. DBSNP=rs206437).

10 THE CLASS ONTOLOGY

As discussed in Section 9.1, variant calls overlapping with transcripts can be annotated with the 'CLASS' flag which refers to a simple ontology describing different types of variants. This ontology is described here in detail.

The values of the CLASS flag are described in the table below:

CLASS	Description
SG	Stop-gain variant caused by base substitution or inframe indel.
ESS	Any variant that alters essential splice-site base (+1, +2, -1, -2).
SS5	Any variant that alters the +5 splice-site base but not an ESS base.
SS	Any variant that alters splice-site base within the first @ssrange intronic bases flanking exon (i.e. +@ssrange to -@ssrange) but not an ESS or SS5 base.

EE	Variant that alters the first or last 3 bases of an exon (i.e. the <i>exon end</i>), but not the frame of the coding sequence.
FS	Frameshifting insertion and/or deletion. It alters length and frame of coding sequence.
IM	Variant that alters initiating methionine start codon.
SL	Variant that causes a stop-loss (i.e. the stop codon is altered).
IF	Inframe insertion and/or deletion. It alters length but not frame of coding sequence.
NSY	Nonsynonymous variant. It alters amino acid(s) but not coding sequence length.
SY	Synonymous variant. It does not alter amino acid or coding sequence length.
INT	Any variant in an intron that does not alter splice-site bases.
5PU	Any variant in 5' untranslated region
3PU	Any variant in 3' untranslated region

Notes:

- A variant can only have one CAVA class. If a variant could potentially be included in more than one class the first class in the list is assigned. For example, a frameshifting deletion that alters the start codon would be CAVA class FS (not IM).
- Nonsynonymous is also known as missense. Stop-gain is also known as nonsense.
- NSY, SY and SG variants include complex variants that delete some bases from the coding sequence and insert different bases of an equal number. As a result, the length of coding sequence is not altered, therefore these variants are not the IF or FS classes.
- Indel variants that affect an exon-intron boundary are always classified as ESS variants.
- The @ssrange option in the configuration file can be used to define the size of the splice site region: i.e. the number of bases into the intron used as splice site in CLASS annotation. The parameter affects SS/INT boundary. The default value of @ssrange is 8 in order to make CLASS and SO ontologies comparable.
- Duplications overlapping the boundary of the 3' splice site region of an intron are classified as INT as they do not alter splice site region: e.g. c.10-12dupC or c.10-13_10-11dupTGC (if @ssrange=12).

The values of the CLASS flag have corresponding SO terms, with some noted exceptions, presented in the table below:

CLASS	SO term	Exception
SG	stop_gained	

ESS	splice_acceptor_variant, splice_donor_variant	The SO term is provided as appropriate for variants altering splice acceptor and splice donor sites, however the same CLASS value is returned for both types of variants.
SS5	splice_donor_5th_base_variant	
SS	intron_variant splice_region_variant	The SO term is always provided up to 8 bases into the intron to match the SO definition, however the CLASS value is provided according to @ssrange, so this will not be a one-to-one mapping when @ssrange value is not 8.
EE	splice_region_variant inframe_deletion, splice_region_variant inframe_insertion, splice_region_variant synonymous_variant, splice_region_variant missense_variant	The SO term is provided as appropriate for inframe deletions and insertions and base substitutions which affect the first and last three bases of the exon, however the same CLASS value is returned for all of these types of variants.
FS	frameshift_variant, splice_region_variant frameshift_variant	The SO term is provided as appropriate for frameshifting indels which do and do not affect the first and last three bases of the exon, however the same CLASS value is returned for both types of variants.
IM	initiator_codon_variant	
SL	stop_lost	
IF	inframe_deletion, inframe_insertion	The SO term is provided as appropriate for inframe deletions and insertions, however the same CLASS value is returned for both types of variants.
NSY	missense_variant	
SY	synonymous_variant	

INT	intron_variant	The SO term is always provided beyond 8 bases into the intron to match the SO definition, however the CLASS value is provided according to @ssrange, so this will not be a one-to-one mapping when @ssrange is not 8.
5PU	5_prime_UTR_variant	
3PU	3_prime_UTR_variant	

11 OUTPUT

Variant annotations discussed in Section 9 can be outputted by CAVA in two different file formats; VCF and TSV. The output format is specified in the configuration file by option flag @outputformat. The structure of output files are described in this section.

11.1 Output in VCF format

If VCF output format is used, variant call annotations are added to the INFO field of the VCF file. Following the VCF specification, annotation flags are separated by semicolons in the INFO field. For example:

```
TYPE=Substitution;TRANSCRIPT=ENST00000379410;GENE=PLEKHN1;TRINFO=+/8.6kb/16/2.4kb;LOC=3UTR;CSN=c.+483A>G;PROTPOS=.;PROTREF=.;PROTALT=.;CLASS=3PU;ALTFLAG=None;DBSNP=rs668558
```

Optionally, if configuration flag @prefix is switched on in the configuration file, variant annotation flags are appended to the original INFO field values of each variant prefixed with “CAVA_”. The ID, CHROM, POS, REF and ALT fields in the output VCF file will be the same as the ID, chromosome, position, reference and alternative allele values in the input file; i.e. even if indels are left or right aligned in some transcript, the original genomic coordinates and alleles are given in the output. The QUAL and FILTER field values are also copied from the input file. Furthermore, if the input VCF file contains the FORMAT column and sample-specific genotype calls, these are also present in the output file.

11.1.1 Multiallelic VCF records

Multiallelic variant calls represented in a single VCF record in the input are also outputted in the same VCF record in the output. Each alternative allele has different TYPE, LOC, CSN, PROTPOS, PROTREF, PROTALT, CLASS, SO, ALTFLAG, ALTANN, ALTCLASS, ALTISO and DBSNP flag values. If reported in a single VCF record, the multiple annotations corresponding to the different alternative alleles are comma-separated in the TYPE, LOC,

CSN, PROTPOS, PROTREF, PROTALT, CLASS, SO, ALTFLAG, ALTANN, ALTCLASS, ALTISO and DBSNP flags. For example, for the following multiallelic VCF record describing a 1-base deletion and a 1-base insertion

#CHROM	POS	ID	REF	ALT	...
8	3443799	.	GA	G,GAA	...

the TYPE and CSN annotation flags are given as follows:

TYPE=Deletion,Insertion;
CSN=c.1098-18delT,c.1098-18dupT;

11.1.2 Variants overlapping with multiple transcripts

Variants that overlap with multiple transcripts have different TRANSCRIPT, GENE, GENEID, TRINFO, LOC, CSN, PROTPOS, PROTREF, PROTALT, CLASS, SO, ALTFLAG, ALTANN, ALTCLASS and ALTISO annotation flag values corresponding to the different transcripts. In this case, the multiple annotations values are colon-separated in these annotation flags. For instance, the following substitution overlaps with two transcripts in the whole exome database, ENST00000438763 and ENST00000452392.

#CHROM	POS	ID	REF	ALT	...
6	32784783	.	C	T	...

The TRANSCRIPT, CSN and CLASS annotations corresponding to the two transcripts are given as follows:

TRANSCRIPT=ENST00000438763:ENST00000452392;
CSN=c.-55G>A:c.1933-54G>A;
CLASS=5PU:INT;

11.1.3 Multiallelic calls overlapping with multiple transcripts

If a VCF record describes a multiallelic variant call that overlaps with multiple transcripts, the two rules above are combined: annotation values referring to different transcripts are colon-separated and values corresponding to different alternative alleles are comma-separated. For example, the following multiallelic substitution overlaps with two transcripts in the whole exome set, ENST00000325203 and ENST00000344683.

#CHROM	POS	ID	REF	ALT	...
8	6389889	.	C	G,A	...

The GENE, LOC and CSN annotation flags are given as follows:

GENE=ANGPT2:MCPH1,ANGPT2:MCPH1;
LOC=Ex2:In12/13,Ex2:In12/13;
CSN=c.408G>C_p.=:c.2214+32439C>G,c.408G>T_p.=:c.2214+32439C>A

Note that both alternative alleles in both genes (ANGPT2 and MCPH1) cause synonymous and intronic changes, respectively.

11.2 Output in TSV format

If TSV output format is used, variant call annotations are written to a TAB-delimited file that by default contains the following 24 columns:

1. ID (i.e. Variant call ID taken from the input file)
2. CHROM (i.e. chromosome of variant)
3. POS (i.e. genomic position of variant)
4. REF (i.e. reference allele of variant)
5. ALT (i.e. alternative allele of variant)
6. QUAL (i.e. QUAL value in the input VCF record)
7. FILTER (i.e. FILTER value in the input VCF record)
8. TYPE (i.e. value of TYPE annotation flag)
9. TRANSCRIPT (i.e. value of TRANSCRIPT annotation flag)
10. GENE (i.e. value of GENE annotation flag)
11. GENEID (i.e. value of GENEID annotation flag)
12. TRINFO (i.e. value of TRINFO annotation flag)
13. LOC (i.e. value of LOC annotation flag)
14. CSN (i.e. value of CSN annotation flag)
15. PROTPOS (i.e. value of PROTPOS annotation flag)
16. PROTREF (i.e. value of PROTREF annotation flag)
17. PROTALT (i.e. value of PROTALT annotation flag)
18. CLASS (i.e. value of CLASS annotation flag)
19. SO (i.e. value of SO annotation flag)
20. IMPACT (i.e. value of IMPACT annotation flag)
21. ALTANN (i.e. value of ALTANN annotation flag)
22. ALTCLASS (i.e. value of ALTCLASS annotation flag)
23. ALTSO (i.e. value of ALTSO annotation flag)
24. DBSNP (i.e. value of DBSNP annotation flag)

11.2.1 Multiallelic calls and/or multiple transcripts

Every line of the output TSV file represents a single variant call. Unlike in the VCF format, annotation information for multiallelic variant calls are split into multiple records. If a variant call overlaps with multiple transcripts, the information is also split into multiple records. Two examples are shown below.

For the single VCF record in Section 11.1.1, two lines will be added to the output TSV file:

```
. 8 3443799 GA G 42 PASS Deletion ENST00000537824 CSMD1 ...
. 8 3443799 GA GAA 42 PASS Insertion ENST00000537824 CSMD1 ...
```

Another example is the single VCF record in Section 11.1.3 which will be represented by four lines in the output TSV file:

```
. 8 6389889 C G 200 PASS Substitution ENST00000325203 ANGPT2 ...
. 8 6389889 C G 200 PASS Substitution ENST00000344683 MCPH1 ...
. 8 6389889 C A 200 PASS Substitution ENST00000325203 ANGPT2 ...
. 8 6389889 C A 200 PASS Substitution ENST00000344683 MCPH1 ...
```

12 VARIANT CALL FILTERING

CAVA offers a number of different options to filter the input variant set so that only the selected subset of calls are annotated and written to the output file.

12.1 Filtering by variant type

Specified in the configuration file by flag `@type`, this option can be used to select only a particular variant type: substitutions, insertions, deletions, complex or all indels. Alternatively, setting the value to 'ALL' makes CAVA annotate and output every variant call regardless of their type.

12.2 Filtering by VCF filter

CAVA can filter out variant call records based on the FILTER field of the input VCF file. If this option is switched on in the configuration file by the boolean flag `@filter`, CAVA will only annotate VCF records that have a PASS filter value.

12.3 Filtering by BED file

CAVA can apply a filter on variant calls based on a compressed BED file describing a set of genomic regions of interest. If this option is used, the program will only annotate and output variants which overlap with any of the genomic regions specified in the BED file. The BED file should use 0-based genomic coordinates. It must be position sorted, compressed by bgzip and indexed by Tabix (see <http://www.htslib.org/doc/tabix.html>)

Both the bgzipped BED file and the index file (e.g. `panel.bed.gz` and `panel.bed.gz.tbi`) should be available in the same directory. The name of the gzipped BED file should be given in the configuration file at option flag `@target`.

12.4 Filtering by gene list

Another option CAVA offers for filtering is restricting variant annotations to a subset of genes of interest. In this case, only variant calls overlapping with particular genes are annotated and outputted. The list of gene (HGNC) symbols of interest should be given in a simple txt file (see an example below). The name of the txt file must be specified in the configuration file by option flag `@genelist`.

Example gene list txt file:

```
BRCA1  
BRCA2  
FANCD2  
...
```

12.5 Filtering by transcript list

Similarly to filtering by gene list, CAVA can restrict variant annotations to a subset of transcripts of interest. In this case, only variant calls overlapping with particular transcripts are annotated and outputted. The list of Ensembl transcript IDs (ENST) should be given in a simple txt file (see an example below). The name of the txt file must be specified in the configuration file by option flag `@transcriptlist`.

Example transcript list txt file:

ENST00000380152

ENST00000358533

ENST00000338591

...

12.6 Filtering by SNP list

CAVA can filter variant annotations based on a list of selected SNPs. In this case, only base substitutions annotated with particular dbSNP identifiers are outputted. The list of dbSNP IDs should be given in a simple txt file (see an example below). The name of the txt file must be specified in the configuration file by option flag `@snplist`.

Example SNP list txt file:

rs4104967

rs206437

rs149472673

...

Note that if more than one of `@genelist`, `@transcriptlist` and `@snplist` are specified, the AND operation is applied: i.e. in order to be outputted variants have to satisfy both or all three conditions.

12.7 Filtering non-annotated calls

Some variants may neither overlap with any of the transcripts in the local Ensembl transcript dataset nor correspond to any SNPs in the local dbSNP dataset. These non-annotated variants get only one non-empty annotation flag; the TYPE flag. CAVA offers an option to filter out these non-annotated calls. If the option is switched on in the configuration file by option flag `@nonannot`, only variant calls with either non-empty TRANSCRIPT or DBSNP annotation flags are outputted.

13 THE ENSEMBL_DB TOOL

As discussed above, CAVA relies on a local transcript database file to perform transcript-based variant annotations. The transcript database file used by CAVA has a specific format and is compressed and indexed with Tabix. A simple tool, `ensembl_db`, is provided in the

CAVA package to generate the correct database file for the Ensembl release of interest (release versions ≥ 70 and version 65 are supported). The transcript database created by `ensembl_db` is restricted to genes on chromosomes 1-22, X, Y, and MT.

For example, a transcript database for Ensembl release 70 can be generated with the following command:

CAVA-1.2.3/ensembl_db -e 70 -o output

where option `-e` specify the Ensembl release version and option `-o` specifies the output file prefix.

This will generate three output files: `output.gz` (the compressed and Tabix-indexed transcript database file), `output.gz.tbi` (the Tabix index file) and `output.txt` (a txt file providing a simple list of included transcripts).

Note that the above command generates a database that contains all protein-coding transcripts of protein-coding genes in the given Ensembl release that have both start and stop codon annotations and are complete.

Running CAVA, one can refer to this database in the configuration file by option flag `@ensembl` (see Section 4):

`@ensembl = output.gz`

13.1 Automatic transcript selection

By default, `ensembl_db` creates a database including all transcripts from an Ensembl release that satisfy the above requirements. Alternatively, one can use the tool's in-built transcript selection pipeline that selects a single transcript for each gene. Automatic transcript selection can be switched on with the `-s` command line flag:

CAVA-1.2.3/ensembl_db -e 70 -o output -s

The resulting database file, `output.gz`, will contain only the selected transcripts. Detailed description of the transcript selection pipeline is given in Appendix 2.

13.2 Including only a custom list of transcripts

Optionally, `ensembl_db` can generate a database for a custom list of transcripts. To this end, the identifiers of transcripts of interest have to be supplied in a simple text file with each ENST identifier given in different line.

For example, let the `input.txt` file contain the following lines:

ENST00000398334

ENST00000263121
ENST00000314074
...

The above file can be inputted to `ensembl_db` with option `-i`:

CAVA-1.2.3/ensembl_db -i input.txt -e 70 -o output

The resulting `output.gz` database will contain only the transcripts from `input.txt` that are present in the specified Ensembl release.

14 DEFAULT TRANSCRIPT DATABASE

Although the `ensembl_db` tool gives the user flexibility to create any transcript database of interest, a default transcript database is also provided as part of the CAVA package. The default transcript database was created by `ensembl_db` for Ensembl release 75 using automatic transcript selection with the following command:

CAVA-1.2.3/ensembl_db -e 75 -s -o ensembl75s

As transcript selection was switched on, the database contains a single transcript for each included gene.

The database files (`.gz` and `.tbi`) and the `.txt` file listing the included transcripts are found in the `CAVA-1.2.3/defaultdb` directory.

Note that if the `@ensembl` flag is not specified in the configuration file, CAVA will automatically use its default transcript database.

15 THE DBSNP_DB SCRIPT

The `dbsnp_db` script is included in the CAVA package and can be used to generate the correct dbSNP database file required by CAVA for dbSNP annotation.

`dbsnp_db` requires the appropriate compressed VCF file (*00-All.vcf.gz*) released by NCBI containing the entire set of SNPs in the dbSNP database. This file can be downloaded from the official dbSNP FTP site: <ftp://ftp.ncbi.nih.gov/snp/organisms/>

For example, the latest dbSNP release can be downloaded as:
ftp://ftp.ncbi.nih.gov/snp/organisms/human_9606/VCF/00-All.vcf.gz

The entire dataset of dbSNP release 146 corresponding to the GRCh37p13 genome build can be downloaded as:

ftp://ftp.ncbi.nih.gov/snp/organisms/human_9606_b146_GRCh37p13/VCF/00-All.vcf.gz

dbSNP_db converts the downloaded 00-All.vcf.gz file into a SNP database file required by CAVA. The user can specify the release of interest that can be lower or equal to the release version (build) of the input 00-All.vcf.gz file. In order to create the dbSNP database file for CAVA, one can run the following command (shown for dbSNP release 137):

```
CAVA-1.2.3/dbSNP_db -d 00-All.vcf.gz -s 137 -o output
```

where option -d specifies the path to the downloaded 00-All.vcf.gz file, -s specifies the requested dbSNP release and -o defines the output file name prefix.

Running the above command, dbSNP_db will create an output file named *output.gz*, which is the compressed and Tabix-indexed database file (an index file *output.gz.tbi* is also created). Running CAVA, one can refer to this database in the configuration file by option flag @dbSNP (see Section 4):

```
@dbSNP = output.gz
```

15.1 Including only a custom list of SNPs

Optionally, dbSNP_db can also generate a database for a custom list of SNPs. To this end, the identifiers of the SNPs of interest have to be supplied in a simple text file with each dbSNP (rs) identifier given in different line.

For example, let the input.txt file contain the following lines:

```
rs4987117
rs206076
rs11571662
...
```

The above file can be inputted to dbSNP_db with option -i:

```
CAVA-1.2.3/dbSNP_db -d 00-All.vcf.gz -s 137 -i input.txt -o output
```

The resulting output.gz database will contain only the SNPs from input.txt that are present in the specified dbSNP release.

16 RUN TIMES

Measured on a 2.9 GHz Intel Core i7 machine, CAVA v1.2.3 was able to annotate 33098 VCF records per minute (approx. 2 million records per hour) without multithreading, or 82747 VCF records per minute (approx. 5 million records per hour) with multithreading.

17 LOG FILE

If switched on in the configuration file by option flag `@log`, CAVA creates a log file during the annotation process, writing out status information and possible error messages. The name of the log file is created from the output file prefix defined at the command line (`-o`) to which the `'.log'` file extension is added.

18 EXAMPLES

Three examples are presented in this final section explaining the configuration file, input file(s) and output file for each. Examples 1 and 2 describe the scenario of annotating high quality variants overlapping gene transcripts from exome data, while Example 3 illustrates annotation of variants from a targeted panel.

18.1 Example 1

18.1.1 Configuration file

In the first example, the configuration file (*config.txt*) is as follows (see explanation below):

```
@inputformat = VCF
@outputformat = VCF
@reference = human_g1k_v37.fasta
@dbsnp = .
@nonannot = TRUE
@filter = FALSE
@type = ALL
@target = .
@genelist = .
@transcriptlist = .
@snplist = .
@logfile = FALSE
@givealt = FALSE
@ontology = CLASS
```

Both the input and output formats are set to VCF by option flags `@inputformat` and `@outputformat`, respectively. The reference genome (*human_g1k_v37.fasta*) located in the current directory is specified by option flag `@reference`. The transcript database (`@ensembl`)

is not specified in the configuration file and CAVA will therefore use the default transcript database (see Section 14). SNP-based annotation is switched off as @dbsnp is empty.

Since option flag @nonannot is set to TRUE, even variants that do not overlap with any transcripts will be outputted. No filtering will be performed based on the FILTER field of the input VCF file as @filter is set to FALSE. CAVA will output all types of variants (@type=ALL). Since @target, @genelist, @transcriptlist and @snplist are all set empty, no further variant filtering is carried out. Only the CLASS ontology will be reported for variants (@ontology=CLASS). Alternative most 5' sequence and CLASS annotations are not reported, only the ALTFLAG flag is outputted (@givealt=FALSE). Finally, as option flag @log is set to FALSE, no log file is created.

18.1.2 Input file

The input file (*input.vcf*) is in VCF format containing 3 records, one of which represents a multiallelic variant call:

#CHROM	POS	ID	REF	ALT	QUAL	FILTER	INFO
9	137735017	1	T	TAGGG	200	PASS	ANYFLAG=1
13	32915330	2	GTGGGTAAAGT	G	926	PASS	ANYFLAG=4
3	197566254	3	T	C,A	200	PASS	ANYFLAG=2

Note that these input VCF records already have information in their INFO field ("ANYFLAG").

18.1.3 Running CAVA

CAVA is run by the following command; command line argument -c specifying the configuration file name, argument -i specifying the input file name and -o specifying the prefix of output file name:

CAVA-1.2.3/cava -c config.txt -i input.vcf -o output

18.1.4 Output file

The output file (*output.vcf*) created by CAVA follows VCF format and it contains the same CHROM, POS, ID, REF, ALT, QUAL and FILTER values as the input VCF file. However, the INFO fields of the three VCF records contain additional information about the results of variant annotation.

INFO field of first VCF record:

ANYFLAG=1;TYPE=Insertion;TRANSCRIPT=ENST00000371817;GENE=COL5A1;GENEID=ENSG00000130635;TRINFO=+/203.1kb/66/8.5kb/1838;LOC=3UTR;CSN=c.+870_+873dupGGGA;PROTPOS=.;PROTREF=.;PROTALT=.;CLASS=3PU;IMPACT=3;ALTFLAG=AnnNotClass

Explanation: According to the INFO field, the first VCF record describes a 4-base insertion (duplication) in the gene COL5A1 (ENSG00000130635) affecting the 3' untranslated region of the ENST00000371817 transcript. The transcript is forward-stranded, has a length of 203.1

kb and includes 66 exons which together make up 8.5 kb of coding sequence + UTR and correspond to a 1838 aa protein sequence. The CSN annotation of the variant is 'c.+870_+873dupGGGA'. The insertion is classified as a 3PU (3' UTR) variant with an impact level of 3, and has an alternative indel representation.

INFO field of second VCF record:

ANYFLAG=4;TYPE=Deletion;TRANSCRIPT=ENST00000380152;GENE=BRCA2;GENEID=ENSG00000139618;TRINFO=+/83.7kb/27/10.9kb/3418;LOC=Ex11-In11/12;CSN=c.6839_6841+6del9;PROTPOS=.;PROTREF=.;PROTALT=.;CLASS=ESS;IMPACT=1;ALTFLAG=None

Explanation: According to the INFO field, the second VCF record describes a 9-base deletion in the gene BRCA2 (ENSG0000013961) overlapping the boundary of Exon 11 and Intron 11/12 in the ENST00000380152 transcript. The transcript is forward-stranded, has a length of 83.7 kb and includes 27 exons which together make up 10.9 kb of coding sequence + UTR and correspond to a 3418 aa protein sequence. The CSN annotation of the variant is 'c.6839_6841+6del9'. The deletion is classified as an ESS (essential splice site) variant, because the +1 and +2 bases are deleted with an impact level of 1, and has only one indel representation.

INFO field of third VCF record:

ANYFLAG=2;TYPE=Substitution,Substitution;TRANSCRIPT=ENST00000334859,ENST00000334859;GENE=LRCH3,LRCH3;GENEID=ENSG00000186001,ENSG00000186001;TRINFO=+/80.3kb/19/2.3kb/712,+/80.3kb/19/2.3kb/712;LOC=Ex10,Ex10;CSN=c.1314T>C_p.=,c.1314T>A_p.Tyr438X;PROTPOS=438,438;PROTREF=Y,Y;PROTALT=Y,X;CLASS=SY,SG;IMPACT=3,1;ALTFLAG=None,None

Explanation: According to the INFO field, the third VCF record describes a multiallelic variant with both alternative alleles representing a base substitution. The substitutions are located within the gene LRCH3 (ENSG00000186001), in Exon 10 of the ENST00000334859 transcript. The transcript is forward-stranded, has a length of 80.3 kb and includes 19 exons which together make up 2.3 kb of coding sequence + UTR and corresponds to a 712 aa protein sequence. The two alternative alleles are flagged with different CSN and CLASS annotations. The CSN annotation of the first base substitution is 'c.1314T>C_p.=' and it is classified as an SY (synonymous) variant with an impact level of 3. By contrast, the CSN annotation of the second substitution is 'c.1314T>A_p.Tyr438X' and it is classified as an SG (stop-gain) variant as amino acid Tyr438 changes into a stop codon and has impact level of 1.

18.2 Example 2

18.2.1 Configuration file

In the second example, the configuration file (*config.txt*) is as follows (see explanation below):

```
@inputformat = VCF
@outputformat = TSV
@reference = human_g1k_v37.fasta
@dbsnp = dbSNP138.gz
@nonannot = FALSE
```



```

@filter = TRUE
@type = DELETION
@target = .
@genelist = genes.txt
@transcriptlist = .
@snplist = .
@logfile = TRUE
@givealt = FALSE
@ontology = CLASS

```

The input format is set to VCF, the output format is set to TSV by option flags `@inputformat` and `@outputformat`, respectively. The reference genome (*human_g1k_v37.fasta*) is specified by option flag `@reference`. Both transcript-based and SNP-based annotations are performed (using the default transcript database and a dbSNP138.gz database specified by option flag `@dbSNP`). Unlike in the previous example, non-annotated variants are not written to the output file (`@nonannot=FALSE`). In addition, records with no PASS value in their VCF FILTER field are also filtered out (`@filter=TRUE`). As set by option flag `@type`, only deletions are outputted. Furthermore, a text file containing a list of gene identifiers is provided (specified by option flag `@genelist`). Only variants that overlap with at least one gene on this list are outputted. Only the CLASS ontology will be reported for variants (`@ontology=CLASS`). Alternative most 5' sequence and CLASS annotations are not reported, only the ALTFLAG flag is outputted (`@givealt=FALSE`). Finally, since option flag `@log` is set to TRUE, a log file is being created during variant annotation.

18.2.2 Input file

The input file (*input.vcf*) is in VCF format containing 7 records; 6 deletions and 1 base substitution:

#CHROM	POS	ID	REF	ALT	QUAL	FILTER	INFO
1	152192825	1	C	T	200	PASS	.
1	152195728	2	AT	A	200	PASS	.
5	176070830	3	GC	G	78	PASS	.
5	176071205	4	AG	A	106	PASS	.
7	25266569	5	TTAA	T	200	PASS	.
17	15341736	6	CAG	C	54	SomeFilter	.
17	15670733	7	CA	C	20	PASS	.

18.2.3 Gene list file

Furthermore, let the gene list file (*genes.txt*) contain the following four gene symbols:

```

HRNR
EIF4E1B
BRCA1
BRCA2

```

18.2.4 Running CAVA

The same as in Section 18.1.3.

18.2.5 Output file

The output file (*output.txt*) created by CAVA follows the TAB-separated TSV format containing 3 variants described in the following 21 columns:

ID	CHROM	POS	REF	ALT	QUAL	FILTER	TYPE	TRANSCRIPT	GENE
2	1	152195728	AT	A	200	PASS	Deletion	ENST00000368801	HRNR
3	5	176070830	GC	G	78	PASS	Deletion	ENST00000318682	EIF4E1B
4	5	176071205	AG	A	106	PASS	Deletion	ENST00000318682	EIF4E1B

GENEID	TRINFO	LOC	CSN	PROTPOS	PROTREF	PROTALT
ENSG00000197915	-/12.1kb/3/9.6kb/2850	Ex2	c.1delA_p.Met1?	1	M	C
ENSG00000175766	+/16.0kb/9/2.0kb/242	In5/6	c.296+97delC	.	.	.
ENSG00000175766	+/16.0kb/9/2.0kb/242	In5/6	c.297-170delG	.	.	.

CLASS	IMPACT	ALTFLAG	DBSNP
FS	1	AnnAndClass	.
INT	3	AnnNotClass	.
INT	3	AnnNotClass	.

Out of the seven input VCF records, only three variant calls are outputted. Record 1 is not included in the output because it is a substitution and only deletions are annotated in this example. Record 5 is filtered out because this deletion affects the gene NPVF which is not present in *genes.txt*. Record 6 is excluded because it has a non-PASS FILTER value in the input VCF file. Finally, record 7 does not overlap with any Ensembl transcript so it is filtered out as a non-annotated variant.

The first variant in the output file (record 2) describes a single base deletion within the gene HRNR (ENSG00000197915), in Exon 2 of the ENST00000368801 transcript. The transcript is reverse-stranded, has a length of 12.1 kb and includes 3 exons which together make up 9.6 kb of coding sequence and corresponds to a 2850 aa protein sequence. The CSN annotation of the variant is ‘c.1delA_p.Met1?’. The deletion is classified as an FS (frameshift) variant and has an alternative indel representation which affects its class. The other two outputted variants are also single base deletions, both within the gene EIF4E1B (ENSG00000175766), in Intron 5/6 of the ENST00000318682 transcript. The transcript is forward-stranded, has a length of 16.0 kb and includes 9 exons which together make up 2.0 kb of coding sequence and corresponds to a 242 aa protein sequence. The CSN annotations of the two deletions are ‘c.296+97delC’ and ‘c.297-170delG’, respectively. Both are classified as INT (intronic) variants and have alternative indel representations.

18.2.6 Log file

The following log file (*output.log*) is written during variant annotation:

```
2016-08-19 15:41:40,964 INFO: CAVA v1.2.3 started.
2016-08-19 15:41:40,964 INFO: Configuration file - config.txt
2016-08-19 15:41:40,964 INFO: Input file (VCF) - input.vcf
2016-08-19 15:41:40,964 INFO: Output file (TSV) - output.txt
2016-08-19 15:41:40,965 INFO: Gene list loaded.
2016-08-19 15:41:40,965 INFO: 7 records to be annotated.
2016-08-19 15:41:40,968 INFO: Connected to Ensembl database.
2016-08-19 15:41:40,968 INFO: Connected to dbSNP database.
2016-08-19 15:41:40,971 INFO: Connected to reference genome.
2016-08-19 15:41:40,976 INFO: Process 1 - variant annotation started.
2016-08-19 15:41:40,976 INFO: 10% of records annotated.
2016-08-19 15:41:40,993 INFO: 20% of records annotated.
2016-08-19 15:41:40,997 INFO: 30% of records annotated.
2016-08-19 15:41:41,000 INFO: 40% of records annotated.
2016-08-19 15:41:41,004 INFO: 50% of records annotated.
2016-08-19 15:41:41,005 INFO: 100% of records annotated.
2016-08-19 15:41:41,006 INFO: Output file = 0.5 Kbyte
2016-08-19 15:41:41,006 INFO: CAVA successfully finished
```

18.3 Example 3

18.3.1 Database preparation

The third example illustrates how one can generate and use custom transcript and dbSNP databases. Let the file *transcripts.txt* contain the following Ensembl transcript identifiers of interest:

```
ENST00000327337
ENST00000358821
ENST00000369902
```

The `ensembl_db` tool is used to generate the corresponding transcript database (based on Ensembl release 65):

```
CAVA-1.2.3/ensembl_db -i transcripts.txt -e 65 -o custom_transcripts
```

Similarly, let the file *snps.txt* contain the following dbSNP identifiers of interest:

```
rs11147489
rs4263028
rs111588517
```

The corresponding SNP database (based on dbSNP version 138) is generated with the `dbsnp_db` script:

```
CAVA-1.2.3/dbsnp_db -d 00-All.vcf.gz -s 138 -i snps.txt -o custom_SNPs
```

18.3.2 Configuration file

The configuration file (*config.txt*) is as follows (see explanation below):

```
@inputformat = TXT
@outputformat = VCF
@reference = human_g1k_v37.fasta
@ensembl = custom_transcripts.gz
@dbsnp = custom_SNPs.gz
@nonannot = FALSE
@filter = FALSE
@type = SUBSTITUTION
@target = .
@genelist = .
@transcriptlist = .
@snplist = .
@logfile = FALSE
@givealt = FALSE
@ontology = CLASS
```

The input format is set to TXT, the output format is set to VCF by option flags `@inputformat` and `@outputformat`, respectively. Both transcript-based and SNP-based annotations are performed, using the custom transcript and dbSNP databases generated above (*custom_transcripts.gz* and *custom_SNPs.gz*). As in the previous example, non-annotated variants are not written to output (`@nonannot=FALSE`). Finally, only base substitutions are annotated and outputted (`@type=SUBSTITUTION`).

18.3.3 Input file

The input file (*input.txt*) is in TXT format (see Section 5.2) containing five records, one of which describes a multiallelic call.

#ID	CHROM	POS	REF	ALT
1	12	50745863	C	A,G
2	12	50759433	AAATG	A
3	13	32906980	A	G
4	13	32929478	C	T
5	18	72228124	A	G

18.3.4 Running CAVA

The extension of the input file is changed to TXT:

```
CAVA-1.2.3/cava -c config.txt -i input.txt -o output
```

18.3.5 Output file

The output file (*output.vcf*) created by CAVA follows VCF format and the results of variant annotation are reported in the INFO field. Out of the 5 input records only 3 are present in the output. Record 2 is filtered out as it is not a substitution. Record 3 is excluded because, although it is a substitution affecting BRCA2, this gene is not included in the custom transcript database, therefore the variant is considered non-annotated. Record 4 is written to the output because its dbSNP identifier is included in the custom SNP database. As records 1 and 5 describe substitutions in the FAM186A and CNDP1 genes which are included in the custom transcript database, these variants are written to the output file.

INFO field of Record 1:

TYPE=Substitution,Substitution;TRANSCRIPT=ENST00000327337,ENST00000327337;GENE=FAM186A,FAM186A;GENEID=ENSG00000185958,ENSG00000185958;TRINFO=-/69.2kb/8/7.1kb/2351,-/69.2kb/8/7.1kb/2351;LOC=Ex4,Ex4;CSN=c.4752G>T_p.=,c.4752G>C_p.=;PROTPOS=1584,1584;PROTREF=A,A;PROTALT=A,A;CLASS=SY,SY;IMPACT=3,3;ALTFLAG=None,None;DBSNP=.,.

Explanation: As shown by the INFO field of record 1, this multiallelic substitution is located within the gene FAM186A, in Exon 4 of the ENST00000327337 transcript. The transcript is reverse-stranded, has a length of 69.2 kb and includes 8 exons which together make up 7.1 kb of coding sequence + UTR and corresponds to a 2351 aa protein sequence. The CSN annotations for the different alternative alleles are 'c.4752G>T_p.=' and 'c.4752G>C_p.=' , respectively. Both are classified as SY (synonymous) variants.

INFO field of Record 4:

TYPE=Substitution;TRANSCRIPT=.;GENE=.;GENEID=.;TRINFO=.;LOC=.;CSN=.;PROTPOS=.;PROTREF=.;PROTALT=.;CLASS=.;IMPACT=.;ALTFLAG=.;DBSNP=rs11147489

Explanation: The INFO field of record 4 reports that this substitution does not overlap with any transcripts in the custom database but has a dbSNP identifier (rs11147489).

INFO field of Record 5:

TYPE=Substitution;TRANSCRIPT=ENST00000358821;GENE=CNDP1;GENEID=ENSG00000150656;TRINFO=+/50.6kb/12/2.2kb/507;LOC=Ex4;CSN=c.337A>G_p.Ile113Val;PROTPOS=113;PROTREF=I;PROTALT=V;CLASS=NSY;IMPACT=2;ALTFLAG=None;DBSNP=rs4263028

Explanation: Finally, as shown by the INFO field of record 5, this substitution is located within the gene CNDP1, in Exon 4 of the ENST00000358821 transcript. The transcript is forward-stranded, has a length of 50.6 kb and includes 12 exons which together make up 2.2 kb of coding sequence + UTR and corresponds to a 507 aa protein sequence. The CSN annotation is 'c.337A>G_p.Ile113Val' and the variant is classified as NSY (non-synonymous) variant. It also has a dbSNP identifier reported (rs4263028).

19 CONTACT

Please send any bug reports, comments or feature requests for CAVA to the Rahman team at rahmanlab@icr.ac.uk.

APPENDIX 1:

THE CLINICAL SEQUENCING NOTATION (CSN) v1.0

Background

A fixed, standardized, versioned nomenclature for reporting clinical sequence data, identical for all mutation detection platforms and readily interchangeable with historic data is of vital importance. It allows integration of sequencing data from multiple sources and facilitates more accurate clinical interpretation of genomic information. The Clinical Sequencing Notation (CSN) aims to achieve this. It follows the principles of the existing HGVS nomenclature [1], with minor amendments to ensure compatibility and integration of historical clinical sequence data, whilst also allowing high-volume automated output from NGS platforms.

The aims of CSN

The aims of CSN are

- To provide a fixed, standardized system in which each variant has a single notation
- To use a logical terminology understandable to non-experts
- To provide a nomenclature that allows easy visual discrimination between the major classes of variant in clinical genomics

Description of CSN

The CSN is described in detail below. General points:

- CSN provides nomenclature for three basic types of variant, defined at the nucleotide level
 - Substitution: a change of one base to another
 - Indel: insertion or deletion of one or more bases
 - Complex: a change involving consecutive base other than simple loss or gain of sequence.
- Other types of variation exist (e.g. inversions, conversions, or translocations) but are not currently encompassed by the CSN.
- All CSN descriptions represent a variant in a single allele.
- All CSN variants are reported with a single descriptor. Nucleotide and amino acid level descriptions are joined with an underscore.

Use of the CSN requires the following:

- Reference sequence transcript
- Variant sequence to be annotated

The reference sequence transcript

The reference sequence transcript is the nucleotide (DNA) and amino acid (protein) sequence against which variation in the sequence being annotated is described.

The nucleotide sequence defines the amino acid sequence according to the genetic triplet code: three nucleotides (termed a 'codon') code for one amino acid [2].

The transcript includes:

- Exons – regions of sequence which are translated into the protein. The protein code starts with an initiation codon and ends with a stop codon.
- Introns – regions of sequence between exons that are 'spliced' out and not included in the protein
- UTR – regions that are transcribed but not translated – the region before the protein starts is called the 5' UTR and the region after the protein ends is called the 3' UTR.

Reference sequence position numbering

- The sequence is numbered from the 5' to 3' direction of the transcript.
- Nucleotide (DNA) sequence level positions are prefaced with "c."
- Amino acid (protein) sequence level positions are prefaced with "p."
- "c.1" is always the first protein-coding base, e.g. "A" of the initiation codon.
- "p.1" is always the initiation codon.
- Positions within the intron and UTR are numbered in relation to the nearest exon. Positions 5' of the nearest exon are denoted with "-". Positions 3' of the nearest exon are denoted with "+".
- A base that occupies the central position within an intron comprised of an odd number of bases, i.e. equidistant from neighboring exons, is assigned "+".

Examples:

c.10T	The tenth protein-coding base.
c.-10T	The tenth base position in the sequence situated 5' of the translation initiation codon.
c.+10T	The tenth base position in the sequence situated 3' of the translation termination codon.
c.100+10T	Position c.100 is the final coding base of the exon. The base indicated is the tenth intronic base, situated 3' from the splice donor site.
c.101-10G	Position c.101 is the first coding base of the exon. The base indicated is the tenth intronic base, situated 5' to the splice acceptor site.

Variant description

- The CSN uses the three letter code to describe amino acids [2] with the exception of 'X' which is used to describe any stop codon.
- A single descriptor with the nucleotide and amino acid (if appropriate) level changes, linked with an underscore is given.
- Variants that do not change the amino acid (synonymous) are described at the amino acid level with "p.="
- Variants that change the amino acid (nonsynonymous, also called missense) are described at the amino acid level with the reference amino acid(s), the position, and the variant amino acid(s).
- Stop-gain variants (also called nonsense) are designated at the amino acid level with 'X'.
- Variants which change the initiation codon are described at the amino acid level as with '?', e.g. "p.Met1?"
- Variants which result in loss of the termination codon (stop-loss variants) are described at the amino acid level as "p.extX" followed by either the length of the predicted amino acid extension when a subsequent stop codon exists within the continued reading frame, or "?" when no subsequent stop codon exists within the 3' UTR sequence along the continued reading frame.
- Only nucleotide level information is provided for variants in UTR or intronic sequence.
- Indel variants and complex variants which alter the protein sequence length can either shift the reading frame (frameshifting) or leave the reading frame unaltered (inframe).
- Only a nucleotide level description is given for frameshifting variants. Most lead to mRNA decay and removal of the transcript, rather than a truncated protein [3]
- Indels and complex variants with multiple possible representations according to the reference sequence are described in the CSN at the most 3' position in the coding transcript.

Substitution variants

Substitution variants are those in which one base is ‘substituted’ for another base. They are described at the nucleotide level with the position, followed by the reference base(s), “>”, and the variant base(s) and at the amino acid level according to the impact.

Examples:

c.99A>C_p.=	C is substituted for base 99 (A). There is no change to the amino acid (synonymous variant).
c.99A>C_p.Gln33His	C is substituted for base 99 (A), resulting in histidine replacing glutamine (nonsynonymous variant).
c.78-3G>T	T is substituted for the base positioned three bases 5’ of the first base (78) of the nearest exon. This is an intronic variant.
c.97A>T_p.Lys33X	T is substituted for base 97, resulting in a stop codon replacing lysine with a stop codon (stop-gain variant).
c.999A>C_p.extX?	C is substituted for base 999, resulting in a change in the termination codon and no new stop codon exists within the 3’ UTR sequence along the continued reading frame (stop-loss variant).
c.999A>C_p.extX12	C is substituted for base 999, resulting in a change in the termination codon with a predicted extension by 12 amino acids where the next stop codon exists within the continued reading frame (stop-loss variant).
c.3G>A_p.Met?	A is substituted for base 3, resulting in a change in the initiating methionine codon.

Indel variants

Indel variants can be further classified into three categories: deletions, insertions and duplications. Each category is described below.

Deletions

Deletions are described at the nucleotide level with the position of the deleted bases followed by “del” and either the deleted bases (deletions of 1-4 bases) or the number of deleted bases (deletions of 5 or more bases).

When more than one base is deleted, the positions of the deleted bases are described as the first and last deleted positions separated by an underscore.

Examples:

c.100delG	The single base 100 is deleted causing a frameshift.
c.121+5delG	The base positioned five bases 3' of the final base (121) of the nearest exon is deleted. This is an intronic variant.
c.100_102delTTT_p.Phe34del	Three bases are deleted, resulting in an inframe deletion of phenylalanine.
c.100_108del9_p.Phe34_Met36del	Nine bases are deleted, resulting in an inframe deletion of three amino acids.
c.107_109delTGC_p.Met36_Pro37delinsThr	Three bases are deleted, resulting in threonine replacing methionine and proline.

Insertions

Insertions are described at the nucleotide level with the position of the insertion followed by “ins” and the inserted bases.

The position of the insertion is described by the two flanking bases of the reference sequence separated by an underscore, thus the numbers presented are always consecutive.

Examples:

c.76_77insT	T is inserted between bases 76 and 77 causing a frameshift.
c.77+4_77+5insAA	AA is inserted between the fourth and fifth bases 3' of the final base (77) of the nearest exon. This is an intronic variant.
c.101_102insCTG_p.Phe34_Thr35insCys	Three bases are inserted between bases 101 and 102, resulting in an inframe insertion of cysteine.
c.103_104insGGT_p.Thr35delinsArgSer	Three bases are inserted between bases 103 and 104, resulting in arginine and serine replacing threonine.

Duplications:

Duplications are a specific category of insertion where the inserted base sequence is a duplication of the immediately preceding reference sequence.

- Duplications can comprise any number of bases, but the category does not include an inserted sequence that duplicates a preceding reference sequence motif more than once.
- Duplications are described at the nucleotide level with the positions of the duplicated sequence followed by “dup” and either the duplicated bases (duplications of 1-4 bases) or the number of duplicated bases (duplications of 5 or more bases).
- The positions of the duplicated sequence are defined by reference to the duplicated reference positions, not the point of insertion.
- When more than one base is duplicated, the positions of the duplicated bases are described as the first and last duplicated positions separated by an underscore.

Examples:

c.5dupT	The single base 5 is duplicated causing a frameshift.
c.5_9dup5	The five bases 5-9 are duplicated causing a frameshift.

c.7_12dup6_p.His3_Pro4dup	The six bases 7-12 are duplicated, resulting in the duplication of histidine and proline.
---------------------------	---

Complex variants

Complex variants are those which involve consecutive bases but are not simple deletions or insertions. They are described at the nucleotide level with the positions of the deleted bases followed by “delins” and the inserted nucleotides.

- Variants involving consecutive bases are described as a single complex variant rather than separate multiple events such as a deletion followed by a substitution.
- Complex variants may or may not result in a change in sequence length.
- Complex variants can be frameshifting or inframe variants (complex indel) or result in no change in sequence length (complex substitution).
- The positions of the deleted bases are described by the range of bases that are lost and being replaced.
- When more than one base is deleted, the positions of the deleted bases are described as the first and last deleted positions separated by an underscore.

Examples:

c.2854_2855delinsAT_p.Ala952Met	The two bases 2854-2855 are replaced with AT, resulting in methionine replacing alanine (complex substitution).
c.112_116delinsTG	The five bases 112-116 are replaced with TG resulting in a frameshift (complex indel).
c.100_102delinsGTTAAG_p.Ser34delinsValLys	The three bases 100-102 are replaced with six bases, resulting in valine and lysine replacing serine (complex indel).

Contact

Please send any comments or requests to the Rahman team at rahmanlab@icr.ac.uk.

References

1. den Dunnen, J.T. and S.E. Antonarakis, *Nomenclature for the description of human sequence variations*. Hum Genet, 2001. **109**(1): p. 121-4.

2. Codons and amino acids (Nomenclature for the description of sequence variants: codons and amino acids) <http://www.hgvs.org/mutnomen/codon.html>
3. Chang, Y. F; Imam, J. S and Wilkinson, M. F. (2007). *The nonsense-mediated decay RNA surveillance pathway*. Annual Review of Biochemistry, 2007 **76**: 51–74.

APPENDIX 2:

AUTOMATIC TRANSCRIPT SELECTION PIPELINE

The pipeline used by the `ensembl_db` tool for automatic transcript selection (as discussed in Section 13.1) is described in this section. When requested by the command line option `-s`, `ensembl_db` follows the below pipeline to select a single transcript for each gene.

- Only protein-coding genes are included in the database.
- Protein-coding transcripts that have both start and stop codon annotations and have no CDS incompleteness flags are below referred to as ‘candidate transcripts’.
- Candidate transcripts with a CCDS flag are referred to as ‘CCDS transcripts’.

Transcript selection rule:

- If the gene has CCDS transcripts, the CCDS transcript which encodes the longest protein sequence is selected. If there are multiple CCDS transcripts corresponding to the longest protein sequence, the one with the longest cDNA sequence is selected.
- If the gene does not have CCDS transcripts, the candidate transcript that encodes the longest protein sequence is selected. If there are multiple candidate transcripts corresponding to the longest protein sequence, the one with the longest cDNA sequence is selected.
- If a gene has no candidate transcripts, the gene is not included.
- Exceptions from the above rules are 101 genes for which manually curated transcripts are selected (see the list in `/path/to/cava/ensembl_db/MCG_transcripts.txt`). For one gene (CDKN2A) not one but two manually selected transcripts are included.